

Old Assignment 1 Answer Sketch

Thomas E. Love

2018-01-06

Contents

1	Looking Over the Data Set	1
1.1	Using the <code>describe</code> function	2
1.2	Glimpsing the data's structure	3
2	Fitting a Logistic Regression Model using the <code>glm</code> function	4
2.1	Storing Probabilities and Linear Predictions	5
3	Some Numerical Summaries of the Fitted Probabilities by Treatment Group	6
3.1	Using <code>dplyr</code> and <code>summarise</code>	6
3.2	Using the <code>by</code> command	6
3.3	Using the <code>tableone</code> library	7
4	Plotting the Fitted Probabilities by Treatment Group	8
4.1	A Fancier Boxplot	9
4.2	A Boxplot using <code>ggplot2</code> , with Notches and Means Indicated	10
4.3	A Violin Plot	11
4.4	A DotPlot to compare the probabilities, via <code>ggplot2</code>	12
4.5	A Density Plot, using <code>ggplot2</code>	13
4.6	Another Density Plot, using <code>ggplot2</code>	14
4.7	Our Old Standby - Comparing Distributions via Histograms	15
4.8	A Back-to-Back Histogram	16
5	What About the ROC Curve and C Statistic?	17

1 Looking Over the Data Set

I will start with a quick summary to be sure things are imported properly in the `hw1` data set...

```
summary(hw1)
```

subject	treatment	cov1	cov2
Min. :101.0	Not Treated:95	Min. :27.03	Min. :29.49
1st Qu.:134.5	Treated :40	1st Qu.:41.74	1st Qu.:46.48

Median :228.0		Median :47.70	Median :53.48
Mean :210.2		Mean :48.36	Mean :52.68
3rd Qu.:261.5		3rd Qu.:56.62	3rd Qu.:59.99
Max. :295.0		Max. :73.28	Max. :73.25
cov3	cov4	female	
Min. : 8.00	Min. : 9.00	Min. :0.0000	
1st Qu.:17.00	1st Qu.:17.00	1st Qu.:0.0000	
Median :20.00	Median :20.00	Median :0.0000	
Mean :20.28	Mean :20.06	Mean :0.4296	
3rd Qu.:24.00	3rd Qu.:23.00	3rd Qu.:1.0000	
Max. :33.00	Max. :33.00	Max. :1.0000	

1.1 Using the describe function

Alternatively, we could use the `describe` function, which is part of the `Hmisc` library...

```
library(Hmisc)
describe(hw1)
```

hw1

7 Variables 135 Observations

subject

n	missing	distinct	Info	Mean	Gmd	.05	.10
135	0	135	1	210.2	70.54	107.7	114.4
.25	.50	.75	.90	.95			
134.5	228.0	261.5	281.6	288.3			

lowest : 101 102 103 104 105, highest: 291 292 293 294 295

treatment

n	missing	distinct
135	0	2

Value	Not Treated	Treated
Frequency	95	40
Proportion	0.704	0.296

cov1

n	missing	distinct	Info	Mean	Gmd	.05	.10
135	0	133	1	48.36	11.83	30.70	34.36
.25	.50	.75	.90	.95			
41.75	47.70	56.62	61.84	64.94			

lowest : 27.03 28.30 28.39 28.73 29.60, highest: 66.37 67.30 67.89 68.39 73.28

cov2

n	missing	distinct	Info	Mean	Gmd	.05	.10
135	0	133	1	52.68	11.08	36.81	39.15
.25	.50	.75	.90	.95			
46.48	53.48	59.99	65.39	67.05			

lowest : 29.49 31.93 32.24 34.24 34.25, highest: 68.90 69.55 70.10 72.03 73.25

cov3

n	missing	distinct	Info	Mean	Gmd	.05	.10
135	0	24	0.995	20.28	5.87	12	14
.25	.50	.75	.90	.95			
17	20	24	27	29			

lowest : 8 9 11 12 13, highest: 28 29 30 32 33

cov4

n	missing	distinct	Info	Mean	Gmd	.05	.10
135	0	23	0.994	20.06	4.929	13	15
.25	.50	.75	.90	.95			
17	20	23	26	28			

lowest : 9 11 12 13 14, highest: 28 29 31 32 33

female

n	missing	distinct	Info	Sum	Mean	Gmd
135	0	2	0.735	58	0.4296	0.4938

1.2 Glimpsing the data's structure

Or, perhaps we just want to see the structure of the data and some of the first few values in each variable, in which case, the `str` command would help, or we could use the `dplyr` library's `glimpse` command...

```
library(dplyr)
glimpse(hw1)
```

Observations: 135

Variables: 7

\$ subject <int> 101, 102, 103, 104, 105, 106, 107, 108, 109, 110, 11...

```
$ treatment <fctr> Treated, Treated, Treated, Treated, Treated, Treate...
$ cov1      <dbl> 38.38, 39.07, 67.30, 61.54, 66.37, 57.08, 50.28, 61....
$ cov2      <dbl> 53.70, 48.53, 53.86, 52.18, 55.60, 44.45, 66.83, 68....
$ cov3      <int> 13, 15, 11, 18, 19, 24, 15, 15, 19, 16, 15, 15, 20, ...
$ cov4      <int> 19, 21, 16, 21, 22, 33, 16, 26, 18, 24, 18, 23, 24, ...
$ female    <int> 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0...
```

The two `treatment` options are named “Treated” and “Not Treated”, as opposed to “Treated” and “Untreated”. Anything so that the thing I wanted to evaluate probabilities for (i.e. Treated as compared to Not Treated) came second alphabetically is appealing, because R, by default, treats the first level in a binary categorical variable as unsuccessful and the second level as successful and generally orders levels of binary variables alphabetically.

2 Fitting a Logistic Regression Model using the `glm` function

```
m1 <- glm(treatment == "Treated" ~
          cov1 + cov2 + cov3 + cov4 + female,
          family=binomial(), data=hw1)
```

We are fitting a model to predict the probability of “Treated” here. If we want to see what’s in `m1`, we can type it in, and see what we get, or ask for a summary, and get some additional details.

```
m1
```

```
Call: glm(formula = treatment == "Treated" ~ cov1 + cov2 + cov3 + cov4 +
          female, family = binomial(), data = hw1)
```

Coefficients:

```
(Intercept)      cov1      cov2      cov3      cov4
    0.23905    0.04159   -0.02512   -0.18594    0.06993
    female
    0.79492
```

```
Degrees of Freedom: 134 Total (i.e. Null); 129 Residual
```

```
Null Deviance:      164.1
```

```
Residual Deviance: 135 AIC: 147
```

```
summary(m1)
```

```
Call:
```

```
glm(formula = treatment == "Treated" ~ cov1 + cov2 + cov3 + cov4 +
     female, family = binomial(), data = hw1)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.9216	-0.7418	-0.4914	0.8746	2.4184

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.23905	2.00561	0.119	0.9051
cov1	0.04159	0.02198	1.892	0.0585 .
cov2	-0.02512	0.02303	-1.091	0.2754
cov3	-0.18594	0.04735	-3.927	8.6e-05 ***
cov4	0.06993	0.05053	1.384	0.1664
female	0.79492	0.44205	1.798	0.0721 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 164.08 on 134 degrees of freedom
 Residual deviance: 135.02 on 129 degrees of freedom
 AIC: 147.02

Number of Fisher Scoring iterations: 4

2.1 Storing Probabilities and Linear Predictions

To store the linear predictions (i.e. log odds of the estimated probabilities,) and the estimated probabilities themselves as part of the hw1 data file, I'll use the following commands:

```
hw1$linpred <- m1$linear.predictors
hw1$prob <- m1$fitted.values
```

The remaining tasks in the assignment essentially require you to obtain some numerical (perhaps) and graphical (mandatory) summaries of the estimated probabilities broken down into the two treatment groups.

3 Some Numerical Summaries of the Fitted Probabilities by Treatment Group

3.1 Using dplyr and summarise

The `dplyr` library can be used to compare the `probs` across the two `treatment` groups, along with some piping commands, to create a little data frame of the summaries you're interested in, as follows:

```
hw1 %>%
  group_by(treatment) %>%
  summarise(avg = mean(prob), med = median(prob), sd = sd(prob)) %>%
  arrange(avg)

# A tibble: 2 x 4
  treatment      avg    med    sd
  <fctr>      <dbl> <dbl> <dbl>
1 Not Treated 0.234 0.188 0.177
2 Treated     0.443 0.430 0.207
```

3.2 Using the by command

Some of you may be more familiar with the `by` command - that works, as well...

```
by(hw1$prob, hw1$treatment, describe)

hw1$treatment: Not Treated
dd[x, ]
      n missing distinct      Info      Mean      Gmd      .05      .10
      95      0       95         1  0.2343  0.1891  0.04319  0.05634
      .25      .50      .75      .90      .95
0.10597 0.18805 0.31990 0.47906 0.55551

lowest : 0.01772653 0.02853116 0.03021174 0.03724054 0.04076704
highest: 0.60114577 0.60689182 0.71653717 0.82781846 0.84216789
-----
hw1$treatment: Treated
dd[x, ]
      n missing distinct      Info      Mean      Gmd      .05      .10
      40      0       40         1  0.4435  0.2411  0.1284  0.1543
      .25      .50      .75      .90      .95
0.2972 0.4299 0.6032 0.6858 0.7658

lowest : 0.05370511 0.10947830 0.12944077 0.15246158 0.15452007
```

highest: 0.68436074 0.69878775 0.76469008 0.78683381 0.86569552

3.3 Using the tableone library

Or, you could use the tableone library to produce a summarized Table 1 describing our results...

```
varlist <- c("prob", "linpred", "cov1", "cov2", "cov3", "cov4", "female")
tab1 <- CreateTableOne(data=hw1,
  vars = varlist,
  factorVars = c("female"),
  strata = c("treatment"))
print(tab1)
```

	Stratified by treatment		p	test
	Not Treated	Treated		
n	95	40		
prob (mean (sd))	0.23 (0.18)	0.44 (0.21)	<0.001	
linpred (mean (sd))	-1.44 (1.09)	-0.29 (1.01)	<0.001	
cov1 (mean (sd))	47.07 (10.20)	51.42 (10.01)	0.025	
cov2 (mean (sd))	53.40 (9.79)	50.98 (9.30)	0.187	
cov3 (mean (sd))	21.45 (4.98)	17.50 (4.74)	<0.001	
cov4 (mean (sd))	19.66 (4.17)	21.00 (4.83)	0.107	
female = 1 (%)	38 (40.0)	20 (50.0)	0.378	

You could even use non-parametric tests, and report quartiles for the continuous covariates...

```
print(tab1, nonnorm=c("prob", "linpred", "cov1", "cov2", "cov3", "cov4"))
```

	Stratified by treatment		p	
	Not Treated	Treated		
n	95	40		
prob (median [IQR])	0.19 [0.11, 0.32]	0.43 [0.30, 0.60]	<0.001	
linpred (median [IQR])	-1.46 [-2.13, -0.75]	-0.28 [-0.86, 0.42]	<0.001	
cov1 (median [IQR])	46.42 [39.97, 53.94]	51.07 [43.32, 60.79]	0.046	
cov2 (median [IQR])	54.08 [47.08, 60.36]	51.66 [44.02, 55.49]	0.122	
cov3 (median [IQR])	21.00 [18.00, 24.00]	18.00 [14.75, 21.00]	<0.001	
cov4 (median [IQR])	20.00 [17.00, 23.00]	19.50 [18.00, 24.00]	0.267	
female = 1 (%)	38 (40.0)	20 (50.0)	0.378	
Stratified by treatment				
	test			
n				
prob (median [IQR])	nonnorm			
linpred (median [IQR])	nonnorm			
cov1 (median [IQR])	nonnorm			
cov2 (median [IQR])	nonnorm			

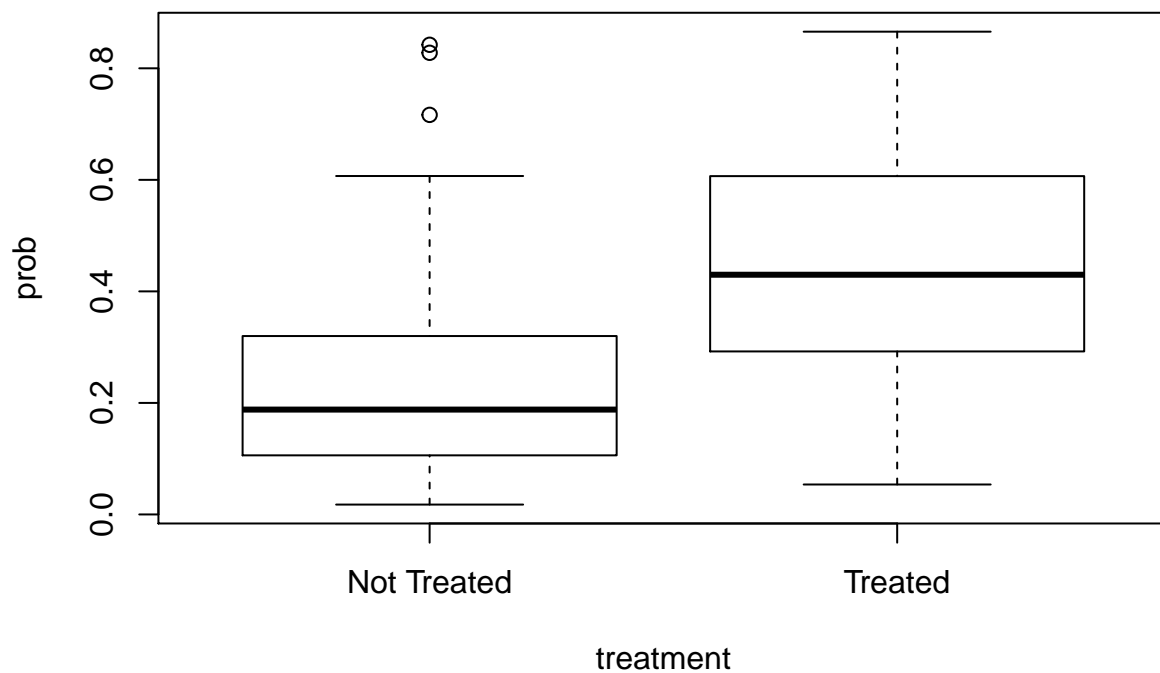
```
cov3 (median [IQR])    nonnorm
cov4 (median [IQR])    nonnorm
female = 1 (%)
```

4 Plotting the Fitted Probabilities by Treatment Group

OK. So we've seen a numerical summary - let's focus on the important issue - a plot.

The default plot in R might help - it's a boxplot.

```
with(hw1,
  plot(prob ~ treatment)
)
```



This could be improved in numerous ways with better titles, etc. but this will get the job done for our purposes in a pinch. As we can see from either the plot or the numerical summaries, we have higher probabilities generally but not inevitably in the **treated** group, and a fair amount of overlap between the two treatment groups in terms of their fitted probabilities.

Having obtained the plot you need, a rational course of action would be to save the results

and your R script or Markdown file, and move on. But I'm not going to stop there.

4.1 A Fancier Boxplot

```
with(hw1,
  boxplot(prob ~ treatment, ylab="Fitted Probability of Treatment",
    main="HW1 Logistic Regression Model Fitted Probabilities",
    horizontal=T, col=c("yellow", "tomato"))
)
```




Tomato is a color? Yes¹.

There are lots and lots of ways to do this sort of thing, some of which involve slicker programming tricks than these. R provides help files for every command, and googling can often help you find the command you need.

¹My favorite reference for R colors is the PDF file at <http://goo.gl/kZqtaI>

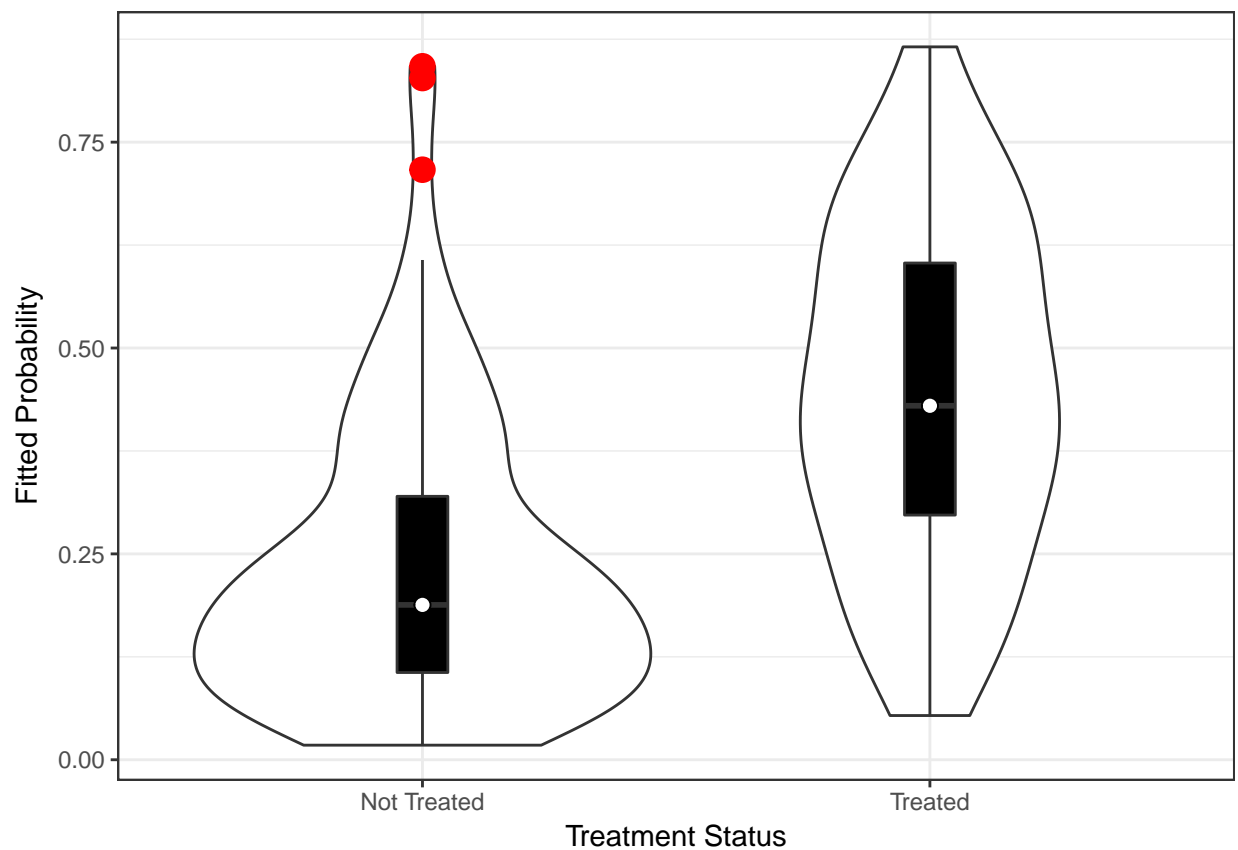
4.2 A Boxplot using ggplot2, with Notches and Means Indicated

```
ggplot(hw1, aes(x = treatment, y = prob)) +  
  geom_boxplot(notch=TRUE) +  
  stat_summary(fun.y="mean", geom="point", shape=23, size = 5, fill = "royalblue")
```




4.3 A Violin Plot

```
ggplot(hw1, aes(x = treatment, y = prob)) +  
  geom_violin() +  
  geom_boxplot(width=.1, fill="black",  
    outlier.colour="red",  
    outlier.size=4) +  
  stat_summary(fun.y=median, geom="point",  
    fill="white", shape=21, size=2.5) +  
  theme_bw() +  
  xlab("Treatment Status") + ylab("Fitted Probability")
```



4.4 A DotPlot to compare the probabilities, via ggplot2


```
ggplot(hw1, aes(x = treatment, y = prob)) +  
  geom_dotplot(binaxis="y", binwidth=0.01, stackdir="center")
```



4.5 A Density Plot, using ggplot2

A possibly more impressive picture would be a density plot. The best way to get this (here, I'll look at the linear probability (i.e. log odds of treatment) results rather than the raw probabilities on a 0-1 scale just to see if we observe something different) uses the ggplot2 library again...

```
ggplot(hw1, aes(x=linpred, fill=treatment)) +  
  geom_density(alpha=0.3) +  
  labs(x="Linear Predictions", colour="Type", shape="Type",  
       title="Linear Predictions By Treatment")
```




One advantage of the linear probabilities over the raw probability estimates is that the log odds results (linear probabilities) are a bit more likely to follow a normalish distribution. Again, it looks like there is fairly substantial overlap in the fitted probabilities across the treatment groups.

4.6 Another Density Plot, using ggplot2

We can use color instead of fill to indicate the densities.


```
ggplot(hw1, aes(x=linpred, color=treatment)) +  
  geom_density() +  
  theme_bw()
```



4.7 Our Old Standby - Comparing Distributions via Histograms

The slickest approach I have here is this:


```
ggplot(hw1, aes(x = prob)) +  
  geom_histogram(fill="white", color="black") +  
  facet_grid(treatment ~ .) +  
  theme(strip.text = element_text(face="bold", size=rel(1.5))) +  
  xlab("Fitted Probability") + ylab("Treatment Status")
```



4.8 A Back-to-Back Histogram

A former student suggested this approach, from the `Hmisc` library. There are likely better ways to get such a plot out of R.

```
# Note: requires Hmisc library be loaded
outhist <- histbackback(split(hw1$prob, hw1$treatment), probability=TRUE,
                        xlim = c(-5,5),
                        main="Back to Back Histogram of Fitted Probabilities")
barplot(-outhist$left, col="light blue", horiz=TRUE, space=0, add=TRUE, axes=FALSE)
barplot(outhist$right, col="pink", horiz=TRUE, space=0, add=TRUE, axes=FALSE)
```



5 What About the ROC Curve and C Statistic?

Recall that our model `m1` was

```
Call: glm(formula = treatment == "Treated" ~ cov1 + cov2 + cov3 + cov4 +  
female, family = binomial(), data = hw1)
```

Since we're looking at a logistic regression, someone in a previous version of this class asked if I could show you how I get the C statistic (area under the ROC curve) for such a model. I usually use the `Epi` library ...

Note that we need to specify the formula (abbreviated form in the ROC function) again, but that's it to get these results.

- The C statistic (area under the curve) for this logistic regression model is 0.786
- Very briefly, the ability of the model's predicted values to discriminate between patients with one outcome vs. the other is quantified by the area under the curve, also called the C statistic or concordance index, which ranges from 0.5 (discrimination is not better than chance) to 1.0 (perfect discriminating power.)
- The ROC procedure comes from signal detection theory and has been adopted into the language of diagnostic testing, essentially treating the response in the logistic regression model as the true status variable, and the set of predictors as the test to be evaluated by things like sensitivity, specificity, and positive and negative predictive values based on dichotomizing along the levels of the predictor set.
- For more on the ROC, visit Wikipedia for Receiver Operating Characteristic. Or try Google.
- A value of 0.786 would indicate a less-than-terrific model in terms of this issue. Values of 0.8 or even 0.9 are usually needed to declare the model to be reasonably accurate in this sense.

```
library(Epi)  
par(mfrow=c(2,1))  
ROC(form=(hw1$treatment=="Treated" ~ hw1$cov1 +  
          hw1$cov2 + hw1$cov3 + hw1$cov4 + hw1$female))
```

